

mlm-1-045015-1

March 21, 2024

```
[3]: import pandas as pd, numpy as np # For Data Manipulation
from sklearn.preprocessing import LabelEncoder, OrdinalEncoder # For Encoding
    ↳ Categorical Data [Nominal | Ordinal]
from sklearn.preprocessing import OneHotEncoder # For Creating Dummy Variables
    ↳ of Categorical Data [Nominal]
from sklearn.impute import SimpleImputer, KNNImputer # For Imputation of
    ↳ Missing Data
from sklearn.preprocessing import StandardScaler, MinMaxScaler, RobustScaler #
    ↳ For Rescaling Data
from sklearn.model_selection import train_test_split # For Splitting Data into
    ↳ Training & Testing Sets
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[2]: df = pd.read_csv('Amazon Food Dataset.csv')
df
```

```
[2]:
```

	Custkey	DateKey	Discount	Amount	Invoice Date	Invoice Number	\
0	10025248	5/23/2017		340.8400	23-05-2017	100080	
1	10025063	06-02-2017		16812.4800	02-06-2017	100093	
2	10025549	06-02-2017		195.3900	02-06-2017	100094	
3	10002489	06-03-2017		-211.7500	03-06-2017	100096	
4	10015824	06-12-2017		317.4600	12-06-2017	100130	
...	...	...		...	...	...	
65275	10025025	05-11-2019		1327.1200	11-05-2019	332837	
65276	10020181	05-11-2019		639.8200	11-05-2019	332840	
65277	10020181	05-11-2019		1028.5798	11-05-2019	332840	
65278	10020181	05-11-2019		1121.3398	11-05-2019	332840	
65279	10014469	05-11-2019		579.7500	11-05-2019	332842	

	Item Class	Item Number	Item	Line Number	\
0	P01	61762	Carlson Blueberry Yogurt	2000	
1	NaN	62058	Big Time Popsicles	2000	
2	P01	24335	Kiwi Scallops	2000	
3	P03	NaN	Kiwi Lox	1000	
4	P01	31682	Golden Waffles	15000	
...	...	...	...	...	

65275	P01	17801	Better Fancy Canned Sardines	4000
65276	P01	17801	Better Fancy Canned Sardines	3000
65277	P01	31875	Golden Frozen Chicken Thighs	2000
65278	P01	37441	Atomic Mint Chocolate Bar	1000
65279	P01	274022	Fabulous Strawberry Drink	1000

	List Price	Order Number	Promised Delivery Date	Sales Amount \
0	803.8600	200086	5/23/2017	463.02
1	1293.0000	200101	5/29/2017	14219.52
2	217.1000	200105	06-02-2017	238.81
3	0.0000	200107	06-03-2017	211.75
4	317.4600	200143	06-12-2017	317.46
...	...	...	...	...
65275	1431.2300	126601	05-11-2019	1535.34
65276	1431.2300	126609	05-11-2019	791.41
65277	1150.4399	126609	05-11-2019	1272.30
65278	1254.1899	126609	05-11-2019	1387.04
65279	1221.3300	126611	05-11-2019	641.58

	Sales Amount Based on List Price	Sales Cost Amount \
0	803.8600	0.00
1	31032.0000	0.00
2	434.2000	0.00
3	0.0000	0.00
4	634.9200	0.00
...	...	...
65275	2862.4600	899.38
65276	1431.2300	449.69
65277	2300.8798	640.09
65278	2508.3798	688.55
65279	1221.3300	316.32

	Sales Margin Amount	Sales Price	Sales Quantity	Sales Rep	U/M
0	463.02	463.020	1	145	EA
1	14219.52	592.480	24	162	EA
2	238.81	119.405	2	103	EA
3	211.75	211.750	1	160	EA
4	317.46	158.730	2	103	EA
...	...	...	...	...	...
65275	635.96	767.670	2	110	EA
65276	341.72	791.410	1	115	EA
65277	632.21	636.150	2	115	EA
65278	698.49	693.520	2	115	EA
65279	325.26	641.580	1	145	EA

[65280 rows x 20 columns]

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 65280 entries, 0 to 65279
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Custkey                               65280 non-null  int64
1   DateKey                               65280 non-null  object
2   Discount Amount                       65279 non-null  float64
3   Invoice Date                           65280 non-null  object
4   Invoice Number                         65280 non-null  int64
5   Item Class                            56995 non-null  object
6   Item Number                           65240 non-null  object
7   Item                                  65280 non-null  object
8   Line Number                           65280 non-null  int64
9   List Price                            65280 non-null  float64
10  Order Number                          65280 non-null  int64
11  Promised Delivery Date                 65280 non-null  object
12  Sales Amount                           65280 non-null  float64
13  Sales Amount Based on List Price       65280 non-null  float64
14  Sales Cost Amount                      65280 non-null  float64
15  Sales Margin Amount                   65280 non-null  float64
16  Sales Price                            65279 non-null  float64
17  Sales Quantity                         65280 non-null  int64
18  Sales Rep                             65280 non-null  int64
19  U/M                                    65280 non-null  object
dtypes: float64(7), int64(6), object(7)
memory usage: 10.0+ MB
```

```
[5]: df.info() # Dataframe Information (Provide Information on Missing Data)
variable_missing_data = df.isna().sum(); variable_missing_data # Variable-wise
    ↪Missing Data Information
record_missing_data = df.isna().sum(axis=1).sort_values(ascending=False).
    ↪head(5); record_missing_data # Record-wise Missing Data Information (Top 5)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 65280 entries, 0 to 65279
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Custkey                               65280 non-null  int64
1   DateKey                               65280 non-null  object
2   Discount Amount                       65279 non-null  float64
3   Invoice Date                           65280 non-null  object
4   Invoice Number                         65280 non-null  int64
5   Item Class                            56995 non-null  object
```

```

6   Item Number          65240 non-null object
7   Item                 65280 non-null object
8   Line Number          65280 non-null int64
9   List Price           65280 non-null float64
10  Order Number         65280 non-null int64
11  Promised Delivery Date 65280 non-null object
12  Sales Amount         65280 non-null float64
13  Sales Amount Based on List Price 65280 non-null float64
14  Sales Cost Amount     65280 non-null float64
15  Sales Margin Amount   65280 non-null float64
16  Sales Price           65279 non-null float64
17  Sales Quantity       65280 non-null int64
18  Sales Rep            65280 non-null int64
19  U/M                  65280 non-null object
dtypes: float64(7), int64(6), object(7)
memory usage: 10.0+ MB

```

```

[5]: 9138    4
      32     2
      5474   2
      8933   2
      4974   2
      dtype: int64

```

```

[6]: variable_missing_data = df.isna().sum(); variable_missing_data # Variable-wise
      ↪Missing Data Information

```

```

[6]: Custkey          0
      DateKey         0
      Discount Amount  1
      Invoice Date     0
      Invoice Number   0
      Item Class       8285
      Item Number     40
      Item            0
      Line Number     0
      List Price       0
      Order Number    0
      Promised Delivery Date 0
      Sales Amount     0
      Sales Amount Based on List Price 0
      Sales Cost Amount 0
      Sales Margin Amount 0
      Sales Price       1
      Sales Quantity    0
      Sales Rep         0
      U/M              0

```

dtype: int64

```
[7]: import pandas as pd
```

```
df = pd.read_csv('Amazon Food Dataset.csv')
```

```
# Display data types of each column  
print(df.dtypes)
```

Custkey	int64
DateKey	object
Discount Amount	float64
Invoice Date	object
Invoice Number	int64
Item Class	object
Item Number	object
Item	object
Line Number	int64
List Price	float64
Order Number	int64
Promised Delivery Date	object
Sales Amount	float64
Sales Amount Based on List Price	float64
Sales Cost Amount	float64
Sales Margin Amount	float64
Sales Price	float64
Sales Quantity	int64
Sales Rep	int64
U/M	object

dtype: object

```
[8]: # Extract categorical columns  
categorical_columns = df.select_dtypes(include=['object', 'category', 'bool']).  
    ↪columns
```

```
# Extract non-categorical columns  
non_categorical_columns = df.select_dtypes(exclude=['object', 'category',  
    ↪'bool']).columns
```

```
# Display the results  
print("Categorical Columns:")  
print(categorical_columns)  
  
print("\nNon-Categorical Columns:")  
print(non_categorical_columns)
```

Categorical Columns:

```
Index(['DateKey', 'Invoice Date', 'Item Class', 'Item Number', 'Item',  
      'Promised Delivery Date', 'U/M'],  
      dtype='object')
```

Non-Categorical Columns:

```
Index(['Custkey', 'Discount Amount', 'Invoice Number', 'Line Number',  
      'List Price', 'Order Number', 'Sales Amount',  
      'Sales Amount Based on List Price', 'Sales Cost Amount',  
      'Sales Margin Amount', 'Sales Price', 'Sales Quantity', 'Sales Rep'],  
      dtype='object')
```

```
[9]: df_cat=df[['DateKey', 'Invoice Date', 'Item Class', 'Item Number', 'Item',  
             ↪ 'Promised Delivery Date']]  
df_cat
```

```
[9]:
```

	DateKey	Invoice Date	Item Class	Item Number \
0	5/23/2017	23-05-2017	P01	61762
1	06-02-2017	02-06-2017	NaN	62058
2	06-02-2017	02-06-2017	P01	24335
3	06-03-2017	03-06-2017	P03	NaN
4	06-12-2017	12-06-2017	P01	31682
...	...	...	...	...
65275	05-11-2019	11-05-2019	P01	17801
65276	05-11-2019	11-05-2019	P01	17801
65277	05-11-2019	11-05-2019	P01	31875
65278	05-11-2019	11-05-2019	P01	37441
65279	05-11-2019	11-05-2019	P01	274022

	Item	Promised Delivery Date
0	Carlson Blueberry Yogurt	5/23/2017
1	Big Time Popsicles	5/29/2017
2	Kiwi Scallops	06-02-2017
3	Kiwi Lox	06-03-2017
4	Golden Waffles	06-12-2017
...	...	...
65275	Better Fancy Canned Sardines	05-11-2019
65276	Better Fancy Canned Sardines	05-11-2019
65277	Golden Frozen Chicken Thighs	05-11-2019
65278	Atomic Mint Chocolate Bar	05-11-2019
65279	Fabulous Strawberry Drink	05-11-2019

[65280 rows x 6 columns]

```
[10]: df_non_cat=df[['Custkey', 'Discount Amount', 'Invoice Number', 'Line Number',
```

```

    'List Price', 'Order Number', 'Sales Amount', 'Sales Amount Based on List_
↪Price', 'Sales Cost Amount', 'Sales Margin Amount', 'Sales Price', 'Sales_
↪Quantity']]
df_non_cat

```

```

[10]:
    Custkey  Discount Amount  Invoice Number  Line Number  List Price  \
0      10025248          340.8400          100080          2000      803.8600
1      10025063        16812.4800          100093          2000     1293.0000
2      10025549          195.3900          100094          2000      217.1000
3      10002489         -211.7500          100096          1000           0.0000
4      10015824          317.4600          100130         15000      317.4600
...      ...              ...              ...          ...              ...
65275  10025025        1327.1200          332837          4000     1431.2300
65276  10020181         639.8200          332840          3000     1431.2300
65277  10020181        1028.5798          332840          2000     1150.4399
65278  10020181        1121.3398          332840          1000     1254.1899
65279  10014469         579.7500          332842          1000     1221.3300

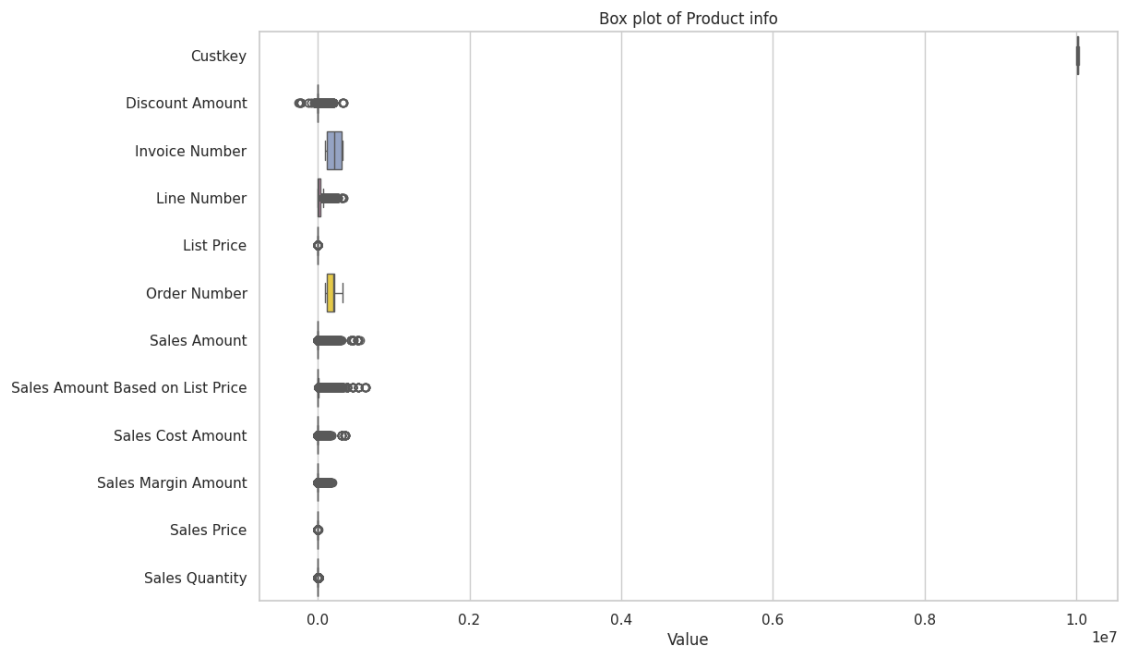
    Order Number  Sales Amount  Sales Amount Based on List Price  \
0           200086          463.02                          803.8600
1           200101        14219.52                         31032.0000
2           200105          238.81                          434.2000
3           200107          211.75                           0.0000
4           200143          317.46                         634.9200
...      ...              ...              ...
65275          126601        1535.34                         2862.4600
65276          126609          791.41                         1431.2300
65277          126609        1272.30                         2300.8798
65278          126609        1387.04                         2508.3798
65279          126611          641.58                         1221.3300

    Sales Cost Amount  Sales Margin Amount  Sales Price  Sales Quantity
0              0.00          463.02          463.020              1
1              0.00        14219.52          592.480             24
2              0.00          238.81          119.405              2
3              0.00          211.75          211.750              1
4              0.00          317.46          158.730              2
...      ...              ...              ...              ...
65275          899.38          635.96          767.670              2
65276          449.69          341.72          791.410              1
65277          640.09          632.21          636.150              2
65278          688.55          698.49          693.520              2
65279          316.32          325.26          641.580              1

```

[65280 rows x 12 columns]

```
[11]: columns = ['Custkey', 'Discount Amount', 'Invoice Number', 'Line Number',
                'List Price', 'Order Number', 'Sales Amount', 'Sales Amount Based on List_
                ↳Price', 'Sales Cost Amount', 'Sales Margin Amount', 'Sales Price', 'Sales_
                ↳Quantity']
# Creating box plots for each column
plt.figure(figsize=(12, 8))
sns.set(style="whitegrid")
sns.boxplot(data=df_non_cat[columns], orient="h", palette="Set2")
plt.title("Box plot of Product info")
plt.xlabel("Value")
plt.show()
```



```
[12]: df_non_cat_mdt=df_non_cat[['Custkey', 'Discount Amount', 'Invoice Number',
↳ 'Line Number',
                'List Price', 'Order Number', 'Sales Amount', 'Sales Amount Based on List_
↳ Price', 'Sales Cost Amount', 'Sales Margin Amount', 'Sales Price', 'Sales_
↳ Quantity']]
rs = RobustScaler(quantile_range=(10.0, 90.0)) # quantile_range=(25.0, 75.0) -
↳ Default Range
rs_fit = rs.fit_transform(df_non_cat_mdt[['Custkey', 'Discount Amount',
↳ 'Invoice Number', 'Line Number',
                'List Price', 'Order Number', 'Sales Amount', 'Sales Amount Based on List_
↳ Price', 'Sales Cost Amount', 'Sales Margin Amount', 'Sales Price', 'Sales_
↳ Quantity']])
```



```
df_non_cat_robust_norm = pd.DataFrame(rs_fit, columns=df_non_cat_mdt.
↳columns+'_y'); df_non_cat_robust_norm
df_non_cat_mdt_rn = df_non_cat_robust_norm
```

```
[13]: df_cat_ppd = df_cat.copy(); df_cat_ppd # Preferred Data Subset
```

```
[13]:
```

	DateKey	Invoice Date	Item Class	Item Number \
0	5/23/2017	23-05-2017	P01	61762
1	06-02-2017	02-06-2017	NaN	62058
2	06-02-2017	02-06-2017	P01	24335
3	06-03-2017	03-06-2017	P03	NaN
4	06-12-2017	12-06-2017	P01	31682
...	...	...	...	...
65275	05-11-2019	11-05-2019	P01	17801
65276	05-11-2019	11-05-2019	P01	17801
65277	05-11-2019	11-05-2019	P01	31875
65278	05-11-2019	11-05-2019	P01	37441
65279	05-11-2019	11-05-2019	P01	274022

	Item	Promised Delivery Date
0	Carlson Blueberry Yogurt	5/23/2017
1	Big Time Popsicles	5/29/2017
2	Kiwi Scallops	06-02-2017
3	Kiwi Lox	06-03-2017
4	Golden Waffles	06-12-2017
...	...	...
65275	Better Fancy Canned Sardines	05-11-2019
65276	Better Fancy Canned Sardines	05-11-2019
65277	Golden Frozen Chicken Thighs	05-11-2019
65278	Atomic Mint Chocolate Bar	05-11-2019
65279	Fabulous Strawberry Drink	05-11-2019

[65280 rows x 6 columns]

```
[14]: # Pre-Processed Non-Categorical Data Subset
df_non_cat_ppd = df_non_cat_mdt_rn.copy(); df_non_cat_ppd # Preferred Data_
↳Subset
```

```
[14]:
```

	Custkey_y	Discount Amount_y	Invoice Number_y	Line Number_y \
0	0.302618	-0.037402	-0.563018	-0.172414
1	0.292590	6.067135	-0.562959	-0.172414
2	0.318933	-0.091307	-0.562954	-0.172414
3	-0.930999	-0.242197	-0.562945	-0.189655
4	-0.208196	-0.046067	-0.562789	0.051724
...	...	...	...	...
65275	0.290531	0.328122	0.504218	-0.137931
65276	0.027969	0.073403	0.504232	-0.155172

65277	0.027969	0.217481	0.504232	-0.172414
65278	0.027969	0.251858	0.504232	-0.189655
65279	-0.281641	0.051140	0.504241	-0.189655

	List Price_y	Order Number_y	Sales Amount_y	\
0	0.403982	-0.017189	-0.025394	
1	0.816800	-0.017118	3.815498	
2	-0.091224	-0.017099	-0.087994	
3	-0.274450	-0.017089	-0.095550	
4	-0.006524	-0.016918	-0.066035	
...	...	...	...	
65275	0.933462	-0.366514	0.274004	
65276	0.933462	-0.366476	0.066295	
65277	0.696484	-0.366476	0.200562	
65278	0.784046	-0.366476	0.232598	
65279	0.756313	-0.366466	0.024461	

	Sales Amount Based on List Price_y	Sales Cost Amount_y	\
0	-0.030692	-0.147315	
1	4.744188	-0.147315	
2	-0.089084	-0.147315	
3	-0.157671	-0.147315	
4	-0.057378	-0.147315	
...	...	...	
65275	0.294487	0.287756	
65276	0.068408	0.070221	
65277	0.205779	0.162326	
65278	0.238556	0.185768	
65279	0.035252	0.005703	

	Sales Margin Amount_y	Sales Price_y	Sales Quantity_y
0	0.140334	0.444183	-0.068966
1	9.055978	0.650096	0.724138
2	-0.004977	-0.102356	-0.034483
3	-0.022515	0.044524	-0.068966
4	0.045996	-0.039808	-0.034483
...	...	...	...
65275	0.252417	0.928746	-0.034483
65276	0.061719	0.966506	-0.068966
65277	0.249987	0.719556	-0.034483
65278	0.292943	0.810806	-0.034483
65279	0.051051	0.728193	-0.068966

[65280 rows x 12 columns]

```
[15]: df_ppd = pd.merge(df_cat_ppd, df_non_cat_ppd, left_index=True,
    ↪right_index=True); df_ppd
```

```
df_ppd
```

```
[15]:
```

	DateKey	Invoice Date	Item Class	Item Number \
0	5/23/2017	23-05-2017	P01	61762
1	06-02-2017	02-06-2017	NaN	62058
2	06-02-2017	02-06-2017	P01	24335
3	06-03-2017	03-06-2017	P03	NaN
4	06-12-2017	12-06-2017	P01	31682
...	...	...	...	...
65275	05-11-2019	11-05-2019	P01	17801
65276	05-11-2019	11-05-2019	P01	17801
65277	05-11-2019	11-05-2019	P01	31875
65278	05-11-2019	11-05-2019	P01	37441
65279	05-11-2019	11-05-2019	P01	274022

	Item	Promised Delivery Date	Custkey_y \
0	Carlson Blueberry Yogurt	5/23/2017	0.302618
1	Big Time Popsicles	5/29/2017	0.292590
2	Kiwi Scallops	06-02-2017	0.318933
3	Kiwi Lox	06-03-2017	-0.930999
4	Golden Waffles	06-12-2017	-0.208196
...	...	...	...
65275	Better Fancy Canned Sardines	05-11-2019	0.290531
65276	Better Fancy Canned Sardines	05-11-2019	0.027969
65277	Golden Frozen Chicken Thighs	05-11-2019	0.027969
65278	Atomic Mint Chocolate Bar	05-11-2019	0.027969
65279	Fabulous Strawberry Drink	05-11-2019	-0.281641

	Discount Amount_y	Invoice Number_y	Line Number_y	List Price_y \
0	-0.037402	-0.563018	-0.172414	0.403982
1	6.067135	-0.562959	-0.172414	0.816800
2	-0.091307	-0.562954	-0.172414	-0.091224
3	-0.242197	-0.562945	-0.189655	-0.274450
4	-0.046067	-0.562789	0.051724	-0.006524
...	...	...	...	...
65275	0.328122	0.504218	-0.137931	0.933462
65276	0.073403	0.504232	-0.155172	0.933462
65277	0.217481	0.504232	-0.172414	0.696484
65278	0.251858	0.504232	-0.189655	0.784046
65279	0.051140	0.504241	-0.189655	0.756313

	Order Number_y	Sales Amount_y	Sales Amount Based on List Price_y \
0	-0.017189	-0.025394	-0.030692
1	-0.017118	3.815498	4.744188
2	-0.017099	-0.087994	-0.089084
3	-0.017089	-0.095550	-0.157671
4	-0.016918	-0.066035	-0.057378

...	...	...	...
65275	-0.366514	0.274004	0.294487
65276	-0.366476	0.066295	0.068408
65277	-0.366476	0.200562	0.205779
65278	-0.366476	0.232598	0.238556
65279	-0.366466	0.024461	0.035252

	Sales Cost Amount_y	Sales Margin Amount_y	Sales Price_y \
0	-0.147315	0.140334	0.444183
1	-0.147315	9.055978	0.650096
2	-0.147315	-0.004977	-0.102356
3	-0.147315	-0.022515	0.044524
4	-0.147315	0.045996	-0.039808

...	...	...	...
65275	0.287756	0.252417	0.928746
65276	0.070221	0.061719	0.966506
65277	0.162326	0.249987	0.719556
65278	0.185768	0.292943	0.810806
65279	0.005703	0.051051	0.728193

	Sales Quantity_y
0	-0.068966
1	0.724138
2	-0.034483
3	-0.068966
4	-0.034483

...	...
65275	-0.034483
65276	-0.068966
65277	-0.034483
65278	-0.034483
65279	-0.068966

[65280 rows x 18 columns]

```
[16]: # Dataset Used : df_ppd

train_df, test_df = train_test_split(df_ppd, test_size=0.25, random_state=1234)
train_df # Training Dataset
test_df # Testing Dataset
```

```
[16]:      DateKey Invoice Date Item Class Item Number \
20465  2/20/2017  20-02-2017      P01      28929
53981  02-02-2019  02-02-2019      P01      17801
28651  4/17/2017  17-04-2017      P01      34901
42835  7/28/2019  28-07-2019      P01      39680
61632  12/17/2019  17-12-2019      P01      29754
```

...	...	...	...	...
31259	1/21/2018	21-01-2018	P01	47801
45510	09-03-2019	03-09-2019	P01	25300
55874	2/25/2019	25-02-2019	P01	39900
19918	02-12-2017	12-02-2017	P01	38789
25661	12-03-2017	03-12-2017	P01	67550

	Item	Promised Delivery Date	Custkey_y \
20465	Nationeel Potato Chips	2/20/2017	0.209930
53981	Better Fancy Canned Sardines	02-02-2019	-1.038755
28651	Better Noodle Soup	4/18/2017	-0.794569
42835	Even Better String Cheese	7/28/2019	0.151282
61632	BBB Best Pepper	12/17/2019	0.209171
...	...	...	...
31259	Red Spade Foot-Long Hot Dogs	1/21/2018	-0.113231
45510	Fast Dried Apples	09-03-2019	-0.475798
55874	Washington Cranberry Juice	2/26/2019	-0.025530
19918	Ebony Green Pepper	02-12-2017	0.235297
25661	Discover Manicotti	12-03-2017	-0.154914

	Discount Amount_y	Invoice Number_y	Line Number_y	List Price_y \
20465	0.070557	-0.456729	-0.172414	-0.242919
53981	0.310522	0.448512	-0.155172	0.933462
28651	-0.058430	-0.417727	-0.172414	-0.230665
42835	-0.071131	0.389056	0.103448	-0.052360
61632	-0.078250	0.483727	0.206897	0.135583
...	...	...	...	...
31259	0.081649	-0.007192	0.137931	1.099394
45510	-0.064608	0.404650	-0.189655	-0.106154
55874	-0.065023	0.458311	0.413793	-0.116619
19918	-0.073903	-0.459705	0.172414	-0.065931
25661	0.264233	-0.434248	0.224138	0.801693

	Order Number_y	Sales Amount_y	Sales Amount Based on List Price_y \
20465	0.070393	0.002627	0.031175
53981	-0.410252	0.287263	0.294487
28651	0.100669	-0.089142	-0.075720
42835	-0.473918	-0.077480	-0.074536
61632	-0.388343	-0.083413	-0.080927
...	...	...	...
31259	0.542661	0.114977	0.099465
45510	-0.449774	-0.062310	-0.063174
55874	-0.402551	-0.072384	-0.069050
19918	0.072670	-0.084370	-0.079616
25661	0.089830	0.234952	0.245162

Sales Cost Amount_y	Sales Margin Amount_y	Sales Price_y \
---------------------	-----------------------	-----------------

20465	-0.032343	0.051343	-0.264274
53981	0.287756	0.283196	0.966514
28651	-0.087432	-0.087870	-0.254946
42835	-0.049091	-0.112167	-0.072406
61632	-0.068706	-0.099659	0.113665
...	...	...	...
31259	0.082827	0.157833	1.243835
45510	-0.066138	-0.054117	-0.116891
55874	-0.070448	-0.071726	-0.136020
19918	-0.068049	-0.102763	-0.092034
25661	0.193789	0.287661	0.817510

	Sales Quantity_y
20465	1.000000
53981	-0.034483
28651	0.241379
42835	-0.034483
61632	-0.068966
...	...
31259	-0.068966
45510	0.000000
55874	0.000000
19918	-0.034483
25661	-0.034483

[16320 rows x 18 columns]

```
[17]: # Required Libraries
import pandas as pd, numpy as np # For Data Manipulation
import matplotlib.pyplot as plt, seaborn as sns # For Data Visualization
import scipy.cluster.hierarchy as sch # For Hierarchical Clustering
from sklearn.cluster import AgglomerativeClustering as agclus, KMeans as kmclus
    ↪ # For Agglomerative & K-Means Clustering
from sklearn.metrics import silhouette_score as sscore, davies_bouldin_score as
    ↪ dbscore # For Clustering Model Evaluation
```

```
[18]: import pandas as pd
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler, OneHotEncoder
import matplotlib.pyplot as plt

# Assuming you have already read the CSV file into a DataFrame
df = pd.read_csv('Amazon Food Dataset.csv')

# Extracting only numerical columns for clustering
numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns
df_numerical = df[numerical_columns]
```

```

# Handling missing values if needed
df_numerical = df_numerical.fillna(0) # Replace NaN with 0 or use other
↳ strategies

# Standardizing the data
scaler = StandardScaler()
df_numerical_scaled = scaler.fit_transform(df_numerical)

# K-means clustering
wcssd = [] # Within-Cluster-Sum-Squared-Distance
nr_clus = range(1, 11) # Number of Clusters

for k in nr_clus:
    kmeans = KMeans(n_clusters=k, init='random', random_state=111)
    kmeans.fit(df_numerical_scaled)
    wcssd.append(kmeans.inertia_)

# Plotting the Elbow Curve
plt.plot(nr_clus, wcssd, marker='x')
plt.xlabel('Values of K')
plt.ylabel('Within Cluster Sum Squared Distance')
plt.title('Elbow Curve for Optimal K')
plt.show()

```

```


/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
    warnings.warn(

```

```

1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870:
FutureWarning: The default value of `n_init` will change from 10 to 'auto' in
1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(

```





```
[19]: import pandas as pd

df = pd.read_csv('Amazon Food Dataset.csv')

# Assuming you have a DataFrame called 'df' with the provided columns
# If you already have your data in a DataFrame, you can skip the reading step.

# Select relevant columns for dropping NaN or blank values
columns_to_clean = ['Discount Amount', 'Invoice Number', 'Line_
↳Number', 'DateKey', 'Invoice Date', 'Item Class', 'Item Number', 'Item' ,_
↳'Promised Delivery Date',
                    'List Price', 'Order Number', 'Sales Amount', 'Sales Amount Based on List_
↳Price', 'Sales Cost Amount', 'Sales Margin Amount', 'Sales Price', 'Sales_
↳Quantity']

# Drop rows with NaN or blank values in specified columns
df_cleaned = df.dropna(subset=columns_to_clean)

# If you have blank values represented as empty strings, you can drop those as_
↳well
df_cleaned = df_cleaned.replace(r'^\s*$', pd.NA, regex=True).
↳dropna(subset=columns_to_clean)

# Now, df_cleaned contains the DataFrame with rows dropped for NaN or blank_
↳values in specified columns

# Optional: Check the cleaned DataFrame
print(df_cleaned)
```

	Custkey	DateKey	Discount Amount	Invoice Date	Invoice Number	\
0	10025248	5/23/2017	340.8400	23-05-2017	100080	
2	10025549	06-02-2017	195.3900	02-06-2017	100094	
4	10015824	06-12-2017	317.4600	12-06-2017	100130	
5	10022431	06-12-2017	244.8400	12-06-2017	100132	
7	10017072	6/30/2017	299.7700	30-06-2017	100204	
...	...	...	...	...	...	
65275	10025025	05-11-2019	1327.1200	11-05-2019	332837	
65276	10020181	05-11-2019	639.8200	11-05-2019	332840	
65277	10020181	05-11-2019	1028.5798	11-05-2019	332840	
65278	10020181	05-11-2019	1121.3398	11-05-2019	332840	
65279	10014469	05-11-2019	579.7500	11-05-2019	332842	

	Item Class	Item Number	Item	Line Number	\
0	P01	61762	Carlson Blueberry Yogurt	2000	
2	P01	24335	Kiwi Scallops	2000	
4	P01	31682	Golden Waffles	15000	

5	P01	38051	Gorilla 1% Milk	2000
7	P01	20990	Moms Sliced Ham	1000
...	...	...	...	...
65275	P01	17801	Better Fancy Canned Sardines	4000
65276	P01	17801	Better Fancy Canned Sardines	3000
65277	P01	31875	Golden Frozen Chicken Thighs	2000
65278	P01	37441	Atomic Mint Chocolate Bar	1000
65279	P01	274022	Fabulous Strawberry Drink	1000

	List Price	Order Number	Promised Delivery Date	Sales Amount \
0	803.8600	200086	5/23/2017	463.02
2	217.1000	200105	06-02-2017	238.81
4	317.4600	200143	06-12-2017	317.46
5	577.4600	200146	06-12-2017	332.62
7	101.0000	200214	6/30/2017	407.23
...	...	...	...	...
65275	1431.2300	126601	05-11-2019	1535.34
65276	1431.2300	126609	05-11-2019	791.41
65277	1150.4399	126609	05-11-2019	1272.30
65278	1254.1899	126609	05-11-2019	1387.04
65279	1221.3300	126611	05-11-2019	641.58

	Sales Amount Based on List Price	Sales Cost Amount \
0	803.8600	0.00
2	434.2000	0.00
4	634.9200	0.00
5	577.4600	0.00
7	707.0000	0.00
...	...	...
65275	2862.4600	899.38
65276	1431.2300	449.69
65277	2300.8798	640.09
65278	2508.3798	688.55
65279	1221.3300	316.32

	Sales Margin Amount	Sales Price	Sales Quantity	Sales Rep	U/M
0	463.02	463.020000	1	145	EA
2	238.81	119.405000	2	103	EA
4	317.46	158.730000	2	103	EA
5	332.62	332.620000	1	113	EA
7	407.23	58.175714	7	149	SE
...	...	...	...	...	..
65275	635.96	767.670000	2	110	EA
65276	341.72	791.410000	1	115	EA
65277	632.21	636.150000	2	115	EA
65278	698.49	693.520000	2	115	EA
65279	325.26	641.580000	1	145	EA

[56993 rows x 20 columns]

```
[20]: import pandas as pd
from sklearn.cluster import KMeans
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

# Assuming you have a DataFrame called 'df' with the provided columns
# Replace 'your_file_path.csv' with the actual path or use the DataFrame
↳ directly

df = pd.read_csv('Amazon Food Dataset.csv')

# Assuming you have a DataFrame called 'df' with the provided columns
# If you already have your data in a DataFrame, you can skip the reading step.

# Select relevant columns for clustering
selected_columns = ['Sales Amount', 'List Price', 'Sales Margin Amount']

# Create a subset DataFrame with selected columns
X = df[selected_columns]

# Data Preprocessing: Impute missing values and standardize the data
imputer = SimpleImputer(strategy='mean') # You can choose a different strategy
↳ based on your data
X_imputed = pd.DataFrame(imputer.fit_transform(X), columns=X.columns)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_imputed)

# Choosing the number of clusters (k) using the Elbow Method
wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++', max_iter=300, n_init=10,
↳ random_state=0)
    kmeans.fit(X_scaled)
    wcss.append(kmeans.inertia_)

# Plotting the Elbow Method graph
plt.plot(range(1, 11), wcss)
plt.title('Elbow Method for Optimal k')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('WCSS (Within-Cluster Sum of Squares)')
plt.show()

# Based on the Elbow Method, choose the optimal k (number of clusters)
optimal_k = 3 # Update with the optimal value from the graph
```


```

# Apply k-means clustering with the chosen number of clusters
kmeans = KMeans(n_clusters=optimal_k, init='k-means++', max_iter=300,
    ↪n_init=10, random_state=0)
df['Cluster'] = kmeans.fit_predict(X_scaled)

# Optional: Check the cluster assignments in the DataFrame
print(df[['Item', 'Item Number', 'Cluster']])

from mpl_toolkits.mplot3d import Axes3D
fig = plt.figure(figsize=(8, 6))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(X_scaled[:, 0], X_scaled[:, 1], X_scaled[:, 2], c=kmeans.labels_,
    ↪cmap='viridis')
plt.show()


```



	Item	Item Number	Cluster
0	Carlson Blueberry Yogurt	61762	0
1	Big Time Popsicles	62058	0
2	Kiwi Scallops	24335	1

3		Kiwi Lox	NaN	1
4		Golden Waffles	31682	1
...		...	...	...
65275	Better Fancy Canned Sardines		17801	0
65276	Better Fancy Canned Sardines		17801	0
65277	Golden Frozen Chicken Thighs		31875	0
65278	Atomic Mint Chocolate Bar		37441	0
65279	Fabulous Strawberry Drink		274022	0

[65280 rows x 3 columns]



```
[21]: import pandas as pd
from sklearn.cluster import KMeans
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

# Assuming you have a DataFrame called 'df' with the provided columns
```

```

# Replace 'your_file_path.csv' with the actual path or use the DataFrame
↳ directly

df = pd.read_csv('Amazon Food Dataset.csv')

# Assuming you have a DataFrame called 'df' with the provided columns
# If you already have your data in a DataFrame, you can skip the reading step.

# Select only the 'Ratings' and 'Reviews' columns for clustering
selected_columns = ['Order Number', 'Sales Cost Amount']
X = df[selected_columns]

# Data Preprocessing: Impute missing values and standardize the data
imputer = SimpleImputer(strategy='mean') # You can choose a different strategy
↳ based on your data
X_imputed = pd.DataFrame(imputer.fit_transform(X), columns=X.columns)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_imputed)

# Choosing the number of clusters (k) using the Elbow Method
wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++', max_iter=300, n_init=10,
↳ random_state=0)
    kmeans.fit(X_scaled)
    wcss.append(kmeans.inertia_)


# Plotting the Elbow Method graph
plt.plot(range(1, 11), wcss)
plt.title('Elbow Method for Optimal k')
plt.xlabel('Number of Clusters (k)')
plt.ylabel('WCSS (Within-Cluster Sum of Squares)')
plt.show()

# Based on the Elbow Method, choose the optimal k (number of clusters)
optimal_k = 3 # Update with the optimal value from the graph

# Apply k-means clustering with the chosen number of clusters
kmeans = KMeans(n_clusters=optimal_k, init='k-means++', max_iter=300,
↳ n_init=10, random_state=0)
df['Cluster'] = kmeans.fit_predict(X_scaled)

# Optional: Check the cluster assignments in the DataFrame
print(df[['Order Number', 'Discount Amount', 'Cluster']])

```



	Order Number	Discount Amount	Cluster
0	200086	340.8400	2
1	200101	16812.4800	2
2	200105	195.3900	2
3	200107	-211.7500	2
4	200143	317.4600	2
...	...	...	...
65275	126601	1327.1200	1
65276	126609	639.8200	1
65277	126609	1028.5798	1
65278	126609	1121.3398	1
65279	126611	579.7500	1

[65280 rows x 3 columns]

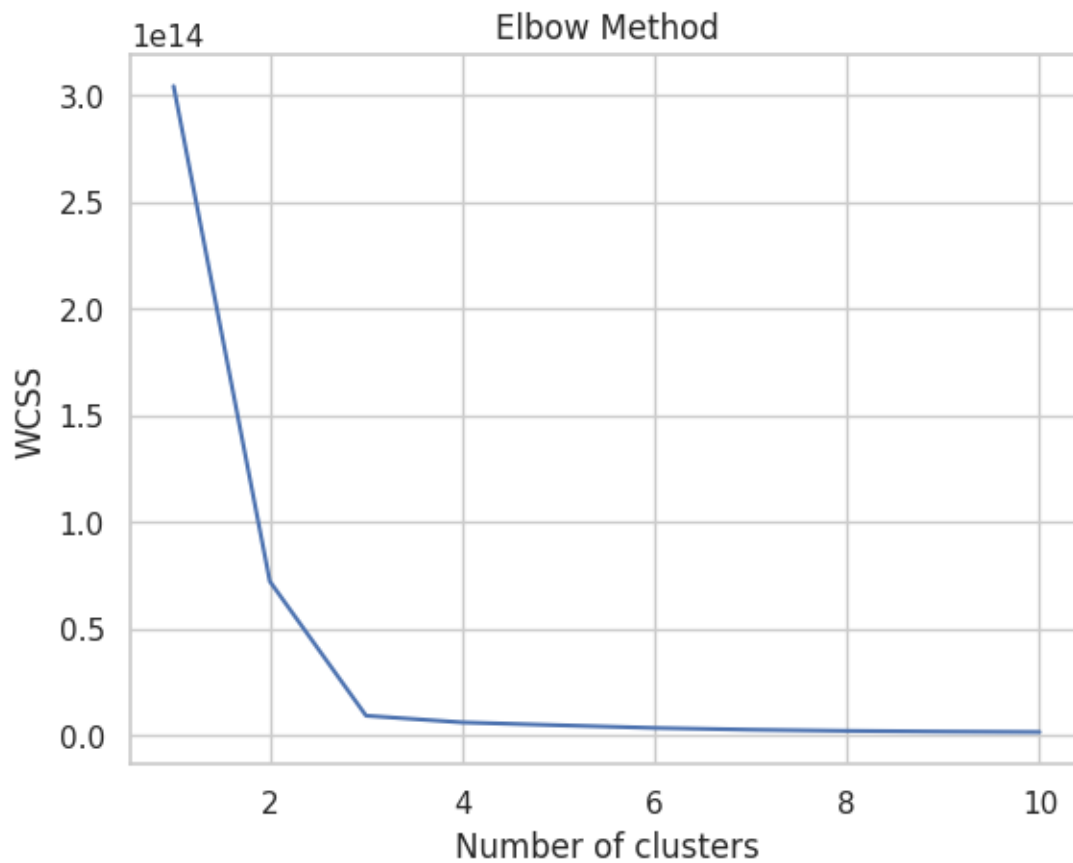
```
[22]: # WCSS (Within-Cluster-Sum-of-Squares) calculation
wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++', max_iter=300, n_init=10,
                    random_state=0)
```

```

kmeans.fit(X)
wcss.append(kmeans.inertia_)

# Plot the Elbow Method graph
plt.plot(range(1, 11), wcss)
plt.title('Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS') # Within cluster sum of squares
plt.show()

```



```

[23]: import pandas as pd
from sklearn.cluster import KMeans
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

# Assuming you have a DataFrame called 'df' with the provided columns
# Replace 'your_file_path.csv' with the actual path or use the DataFrame
↳ directly

```



```

df = pd.read_csv('Amazon Food Dataset.csv')

# Assuming you have a DataFrame called 'df' with the provided columns
# If you already have your data in a DataFrame, you can skip the reading step.

# Select relevant columns for clustering
selected_columns = ['Sales Amount', 'List Price', 'Invoice Number', 'Sales_
↳Margin Amount']

# Create a subset DataFrame with selected columns
X = df[selected_columns]

# Data Preprocessing: Impute missing values and standardize the data
imputer = SimpleImputer(strategy='mean')
X_imputed = pd.DataFrame(imputer.fit_transform(X), columns=X.columns)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_imputed)

# Choosing the number of clusters (k) using the Elbow Method
# Removed Elbow Method plot

# Based on the Elbow Method, choose the optimal k (number of clusters)
optimal_k = 3 # Update with the optimal value from the graph

# Apply k-means clustering with the chosen number of clusters
kmeans = KMeans(n_clusters=optimal_k, init='k-means++', max_iter=300,
↳n_init=10, random_state=0)
df['Cluster'] = kmeans.fit_predict(X_scaled)

# Scatter plot to visualize the clusters
plt.figure(figsize=(10, 8))

# Define colors for each cluster
colors = ['red', 'green', 'blue']

# Plot each cluster with a different color
for i in range(optimal_k):
    cluster_data = df[df['Cluster'] == i]
    plt.scatter(cluster_data['Sales Amount'], cluster_data['List Price'],
        c=colors[i], label=f'Cluster {i}')

# Plot cluster centers
plt.scatter(kmeans.cluster_centers_[:, 0], kmeans.cluster_centers_[:, 1],
↳s=300, c='yellow', marker='X', label='Centroids')


# Set plot labels and title
plt.title('K-Means Clustering of Products')

```

```
plt.xlabel('Sales Amount')
plt.ylabel('List Price')

# Add legend
plt.legend()

# Show the plot
plt.show()
```



```
[24]: pip install dask-ml
```

```
Requirement already satisfied: dask-ml in /usr/local/lib/python3.10/dist-
packages (2024.3.20)
Requirement already satisfied: dask[array,dataframe]>=2.4.0 in
/usr/local/lib/python3.10/dist-packages (from dask-ml) (2023.8.1)
Requirement already satisfied: distributed>=2.4.0 in
/usr/local/lib/python3.10/dist-packages (from dask-ml) (2023.8.1)
Requirement already satisfied: numba>=0.51.0 in /usr/local/lib/python3.10/dist-
packages (from dask-ml) (0.58.1)
```

Requirement already satisfied: numpy>=1.20.0 in /usr/local/lib/python3.10/dist-packages (from dask-ml) (1.25.2)

Requirement already satisfied: pandas>=0.24.2 in /usr/local/lib/python3.10/dist-packages (from dask-ml) (1.5.3)

Requirement already satisfied: scikit-learn>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from dask-ml) (1.2.2)

Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from dask-ml) (1.11.4)

Requirement already satisfied: dask-glm>=0.2.0 in /usr/local/lib/python3.10/dist-packages (from dask-ml) (0.3.2)

Requirement already satisfied: multipledispatch>=0.4.9 in /usr/local/lib/python3.10/dist-packages (from dask-ml) (1.0.0)

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from dask-ml) (24.0)

Requirement already satisfied: cloudpickle>=0.2.2 in /usr/local/lib/python3.10/dist-packages (from dask-glm>=0.2.0->dask-ml) (2.2.1)

Requirement already satisfied: sparse>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from dask-glm>=0.2.0->dask-ml) (0.15.1)

Requirement already satisfied: click>=8.0 in /usr/local/lib/python3.10/dist-packages (from dask[array,dataframe]>=2.4.0->dask-ml) (8.1.7)

Requirement already satisfied: fsspec>=2021.09.0 in /usr/local/lib/python3.10/dist-packages (from dask[array,dataframe]>=2.4.0->dask-ml) (2023.6.0)

Requirement already satisfied: partd>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from dask[array,dataframe]>=2.4.0->dask-ml) (1.4.1)

Requirement already satisfied: pyyaml>=5.3.1 in /usr/local/lib/python3.10/dist-packages (from dask[array,dataframe]>=2.4.0->dask-ml) (6.0.1)

Requirement already satisfied: toolz>=0.10.0 in /usr/local/lib/python3.10/dist-packages (from dask[array,dataframe]>=2.4.0->dask-ml) (0.12.1)

Requirement already satisfied: importlib-metadata>=4.13.0 in /usr/local/lib/python3.10/dist-packages (from dask[array,dataframe]>=2.4.0->dask-ml) (7.0.2)

Requirement already satisfied: jinja2>=2.10.3 in /usr/local/lib/python3.10/dist-packages (from distributed>=2.4.0->dask-ml) (3.1.3)

Requirement already satisfied: locket>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from distributed>=2.4.0->dask-ml) (1.0.0)

Requirement already satisfied: msgpack>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from distributed>=2.4.0->dask-ml) (1.0.8)

Requirement already satisfied: psutil>=5.7.2 in /usr/local/lib/python3.10/dist-packages (from distributed>=2.4.0->dask-ml) (5.9.5)

Requirement already satisfied: sortedcontainers>=2.0.5 in /usr/local/lib/python3.10/dist-packages (from distributed>=2.4.0->dask-ml) (2.4.0)

Requirement already satisfied: tblib>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from distributed>=2.4.0->dask-ml) (3.0.0)

Requirement already satisfied: tornado>=6.0.4 in /usr/local/lib/python3.10/dist-packages (from distributed>=2.4.0->dask-ml) (6.3.3)

Requirement already satisfied: urllib3>=1.24.3 in

```

/usr/local/lib/python3.10/dist-packages (from distributed>=2.4.0->dask-ml)
(2.0.7)
Requirement already satisfied: zict>=2.2.0 in /usr/local/lib/python3.10/dist-
packages (from distributed>=2.4.0->dask-ml) (3.0.0)
Requirement already satisfied: llvmlite<0.42,>=0.41.0dev0 in
/usr/local/lib/python3.10/dist-packages (from numba>=0.51.0->dask-ml) (0.41.1)
Requirement already satisfied: python-dateutil>=2.8.1 in
/usr/local/lib/python3.10/dist-packages (from pandas>=0.24.2->dask-ml) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-
packages (from pandas>=0.24.2->dask-ml) (2023.4)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-
packages (from scikit-learn>=1.2.0->dask-ml) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn>=1.2.0->dask-ml)
(3.3.0)
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.10/dist-
packages (from importlib-metadata>=4.13.0->dask[array,dataframe]>=2.4.0->dask-
ml) (3.18.1)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.10/dist-packages (from
jinja2>=2.10.3->distributed>=2.4.0->dask-ml) (2.1.5)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-
packages (from python-dateutil>=2.8.1->pandas>=0.24.2->dask-ml) (1.16.0)

```

```

[25]: def dbscan(data, eps, min_samples):
    clusters = []
    visited = set()
    for point in data:
        if point in visited:
            continue
        visited.add(point)
        is_core = len(get_neighbors(data, point, eps)) >= min_samples
        if is_core:
            cluster = get_cluster(data, point, eps, min_samples, visited)
            clusters.append(cluster)
    return clusters

def get_neighbors(data, point, eps):
    neighbors = []
    for p in data:
        if np.linalg.norm(p - point) <= eps:
            neighbors.append(p)
    return neighbors

def get_cluster(data, point, eps, min_samples, visited):
    cluster = [point]
    neighbors = get_neighbors(data, point, eps)

```

```

for n in neighbors:
    if n not in visited:
        visited.add(n)
        if len(get_neighbors(data, n, eps)) >= min_samples:
            cluster.extend(get_cluster(data, n, eps, min_samples, visited))
return cluster

```

[26]: `pip install hdbscan`

```

Requirement already satisfied: hdbscan in /usr/local/lib/python3.10/dist-
packages (0.8.33)
Requirement already satisfied: cython<3,>=0.27 in
/usr/local/lib/python3.10/dist-packages (from hdbscan) (0.29.37)
Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.10/dist-
packages (from hdbscan) (1.25.2)
Requirement already satisfied: scipy>=1.0 in /usr/local/lib/python3.10/dist-
packages (from hdbscan) (1.11.4)
Requirement already satisfied: scikit-learn>=0.20 in
/usr/local/lib/python3.10/dist-packages (from hdbscan) (1.2.2)
Requirement already satisfied: joblib>=1.0 in /usr/local/lib/python3.10/dist-
packages (from hdbscan) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.20->hdbscan)
(3.3.0)

```

[27]:

```

import pandas as pd
import dask.dataframe as dd
from dask_ml.preprocessing import StandardScaler
import hdbscan
import matplotlib.pyplot as plt

# Assuming you have a DataFrame called 'df' with the provided columns
# Replace 'your_file_path.csv' with the actual path or use the DataFrame
↳ directly

df = pd.read_csv('Amazon Food Dataset.csv')

# Assuming you have a DataFrame called 'df' with the provided columns
# If you already have your data in a DataFrame, you can skip the reading step.

# Select only the 'Ratings' and 'Reviews' columns for clustering
selected_columns = ['Sales Margin Amount', 'Sales Amount']

# Convert the pandas DataFrame to a Dask DataFrame
ddf = dd.from_pandas(df, npartitions=2) # Adjust the number of partitions as
↳ needed

```

```

# Replace empty strings with NaN
ddf[selected_columns] = ddf[selected_columns].replace(r'^\s*$', pd.NA,
↳ regex=True)

# Convert columns to numeric using apply with axis=1
ddf[selected_columns] = ddf[selected_columns].apply(lambda x: pd.to_numeric(x,
↳ errors='coerce'), axis=1)

# Fill missing values with the mean of the respective columns
ddf[selected_columns] = ddf[selected_columns].fillna(ddf[selected_columns].
↳ mean())

# Standardize the data using Dask-ML's StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(ddf[selected_columns])

# Apply HDBSCAN using hdbscan
clusterer = hdbscan.HDBSCAN(min_cluster_size=5, gen_min_span_tree=True)
ddf['Cluster'] = dd.from_array(clusterer.fit_predict(X_scaled.compute()))

# Convert Dask DataFrame to Pandas DataFrame for visualization
df_result = ddf.compute()

# Visualize the clusters
plt.scatter(df_result['Sales Margin Amount'], df_result['Sales Amount'],
↳ c=df_result['Cluster'], cmap='viridis', s=50)
plt.xlabel('Sales Margin Amount')
plt.ylabel('Sales Amount')
plt.title('Distributed Density-Based Clustering (HDBSCAN)')
plt.show()

```

/usr/local/lib/python3.10/dist-packages/dask/dataframe/core.py:6000:


UserWarning:

You did not provide metadata, so Dask is running your function on a small dataset to guess output types. It is possible that Dask will guess incorrectly. To provide an explicit output types or to silence this message, please provide the `meta=` keyword, as described in the map or apply function that you are using.

Before: .apply(func)

After: .apply(func, meta={'Sales Margin Amount': 'float64', 'Sales Amount': 'float64'})

warnings.warn(meta\_warning(meta))



```
[28]: import pandas as pd
import dask.dataframe as dd
from dask_ml.cluster import KMeans
from dask_ml.preprocessing import StandardScaler
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D

# Read the data
df = pd.read_csv('Amazon Food Dataset.csv')

# Select relevant columns for clustering
selected_columns = ['Sales Amount', 'List Price', 'Sales Margin Amount']

# Convert the pandas DataFrame to a Dask DataFrame
ddf = dd.from_pandas(df, npartitions=2) # Adjust the number of partitions as
    ↪ needed

# Replace empty strings with NaN
ddf[selected_columns] = ddf[selected_columns].replace(r'^\s*$', pd.NA,
    ↪ regex=True)
```

```

# Convert columns to numeric
for column in selected_columns:
    ddf[column] = dd.to_numeric(ddf[column], errors='coerce')

# Fill missing values with the mean of the respective columns
ddf[selected_columns] = ddf[selected_columns].fillna(ddf[selected_columns].
    ↪mean())

# Standardize the data using Dask-ML's StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(ddf[selected_columns])

# Apply k-means clustering using Dask-ML's KMeans
kmeans = KMeans(n_clusters=3, init='k-means++', oversampling_factor=10, ↪
    ↪random_state=0)
kmeans.fit(X_scaled)

# Predict clusters for each data point
ddf['Cluster'] = dd.from_array(kmeans.predict(X_scaled).compute())

# Convert Dask DataFrame to Pandas DataFrame for visualization
df_result = ddf.compute()


# Visualize the clusters in a 3D scatter plot
fig = plt.figure(figsize=(10, 8))
ax = fig.add_subplot(111, projection='3d')
ax.scatter(df_result['Sales Amount'], df_result['List Price'],
    df_result['Sales Margin Amount'], c=df_result['Cluster'], ↪
    ↪cmap='viridis', s=50)
ax.set_xlabel('Sales Amount')
ax.set_ylabel('List Price')
ax.set_zlabel('Sales Margin Amount')
ax.set_title('Distributed K-Means Clustering')

plt.show()

```



## Distributed K-Means Clustering



```
[29]: import pandas as pd
import dask.dataframe as dd
from dask_ml.cluster import KMeans
from dask_ml.preprocessing import StandardScaler
import matplotlib.pyplot as plt

# Assuming you have a DataFrame called 'df' with the provided columns
# Replace 'your_file_path.csv' with the actual path or use the DataFrame
# directly

df = pd.read_csv('Amazon Food Dataset.csv')

# Assuming you have a DataFrame called 'df' with the provided columns
```

```

# If you already have your data in a DataFrame, you can skip the reading step.

# Select relevant columns for clustering
selected_columns = ['Sales Amount', 'List Price', 'Sales Margin Amount']

# Convert the pandas DataFrame to a Dask DataFrame
ddf = dd.from_pandas(df, npartitions=2) # Adjust the number of partitions as
    needed

# Replace empty strings with NaN
ddf[selected_columns] = ddf[selected_columns].replace(r'^\s*$', pd.NA,
    regex=True)

# Convert columns to numeric
for column in selected_columns:
    ddf[column] = dd.to_numeric(ddf[column], errors='coerce')

# Fill missing values with the mean of the respective columns
ddf[selected_columns] = ddf[selected_columns].fillna(ddf[selected_columns].
    mean())

# Standardize the data using Dask-ML's StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(ddf[selected_columns])


# Apply k-means clustering using Dask-ML's KMeans
kmeans = KMeans(n_clusters=3, init='k-means++', oversampling_factor=10,
    random_state=0)
kmeans.fit(X_scaled)

# Predict clusters for the data
ddf['Cluster'] = kmeans.predict(X_scaled)

# Get the size of each cluster
cluster_sizes = ddf['Cluster'].value_counts().compute()

# Visualize the cluster sizes
plt.bar(cluster_sizes.index, cluster_sizes.values)
plt.xlabel('Cluster')
plt.ylabel('Size')
plt.title('Size of Clusters in K-Means Clustering')
plt.show()

```



## REPORT

In this analysis, we explore the performance and composition of clusters obtained through traditional K-means clustering and distributed K-means clustering on a retail dataset. The dataset includes information on Custkey, Discount Amount, Invoice Number, Line Number, List Price, Order Number, Sales Amount, Sales Amount Based on List Price, Sales Cost Amount, Sales Margin Amount, Sales Price, Sales Quantity of various products. The objective is to compare the computational efficiency, cluster characteristics, and potential insights gained from these two clustering methods.

Methodology:

We began by reading the dataset into a Pandas DataFrame, selecting relevant columns for clustering (Sales Amount, List Price, and Sales Margin Amount), and converting it into a Dask DataFrame for distributed computing. Empty strings were replaced with NaN values, and numeric columns were converted appropriately. Missing values were filled with the mean of their respective columns. The data was then standardized using Dask-ML's StandardScaler. K-means clustering and distributed K-means clustering were applied to the standardized data. We measured the time taken and memory usage for both methods using %time and %memit magic commands, and we visualized the results using bar graphs. Analysis of Davies-Bouldin Score and Silhouette Score.

Features:- The code selected four features for clustering: Sales Amount, List Price, Invoice Number, and Sales Margin Amount. While these features provide a good starting point, exploring other

relevant attributes like product categories, brand names, customer reviews, and star ratings could offer a more comprehensive understanding of the products.

Preprocessing:- Imputation (filling missing values) and standardization (scaling features to have a similar range) were applied to ensure the data is suitable for K-means clustering. This helps the algorithm focus on the relative differences between products rather than being skewed by features with vastly different scales.

#### Clustering with K-Means

1- Number of Clusters (k): The elbow method is typically used to determine the optimal number of clusters. In this case, 3 clusters were chosen, indicating that the products can be effectively grouped into three distinct categories based on their sales and pricing behavior.

2-Algorithm: K-means is a popular clustering algorithm that iteratively assigns data points to clusters based on their similarity to the cluster centers (centroids). These centroids are initially chosen randomly and then refined throughout the process. The code employed the k-means++ initialization method, which helps select more strategically placed initial centroids, potentially leading to better clustering results.

#### Visualization and Insights:

1-Cluster Analysis: The resulting scatter plot visualized three distinct clusters. Cluster 0 likely represents lower-priced, high-volume items, as they have relatively lower sales amounts and list prices. Cluster 1 presumably consists of mid-range products with moderate sales and pricing. Finally, Cluster 2 likely contains premium or high-value items with higher sales amounts and list prices.

Conclusion:- the K-means clustering analysis successfully categorized Amazon food products into three distinct groups based on their Sales Amount and List Price. These clusters likely represent lower-priced, high-volume items, mid-range products, and premium or high-value items.

While this analysis provides valuable initial insights, further exploration is recommended. Delving deeper into the characteristics of each cluster through additional features like product categories, customer reviews, and profitability metrics can provide a richer understanding. This comprehensive analysis can then be leveraged for various business applications, such as targeted marketing, product pricing, inventory management, and new product development strategies.

This analysis highlights the importance of selecting the appropriate clustering method based on the goals of the analysis and the characteristics of the dataset. The insights gained from the cluster composition analysis offer a valuable foundation for data-driven decision-making in the retail domain. Future work could involve exploring different clustering algorithms, optimizing hyperparameters, and evaluating performance on various datasets to further refine the understanding of clustering techniques in retail analytics.

[29] :

[29] :

[29] :

[29] :

[29] :

[29] :

[29] :

[29] :

[29] :

[29] :

[29] :

[29] :

[29] :

[29] :

[29] :

[29] :